

Memory Management


- ***Peter Rounce - room G06
p.rounce@cs.ucl.ac.uk***

Memory Protection

Restrict addresses that an application can access:-

- **Limits application to its own code and data**
- **Prevents access to other programs and OS**
- **Prevents access to I/O addresses**

This address restriction is done by Memory Management hardware:-




This Memory Management hardware does address translation as well.

Memory Protection + Address Translation:

Creation of illusion that program has all the memory to itself


Logical Memory

e.g. for Program X



Program, users & CPU's
view of memory
used by Program X

Physical Memory of system



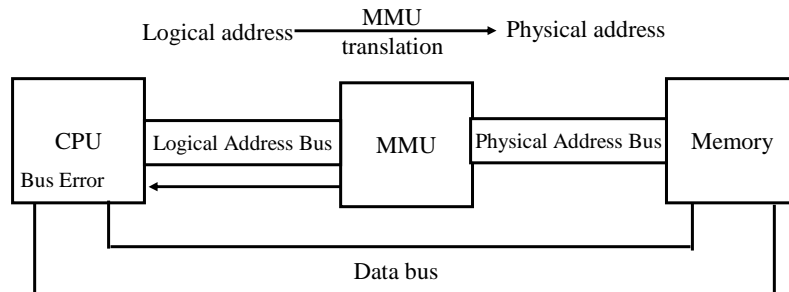
Physical memory is shared
with other programs/processes
and with Operating System

mapped by
address translation

Actual memory
used by Program

Solution: have logical addresses used by program and CPU ; have physical addresses to access memory
Use Address translation to turn logical address into physical address.

Address Translation: basis of Memory Protection and Memory Management



Logical address Address programmed into code and used by CPU
 CPU outputs Logical address on to CPU address Bus
 CPU only knows about Logical addresses

Logical address where CPU considers data/instruction to be stored
 Physical address address where data/instruction is really stored

Note: Physical addresses are never stored in the program, only appear on address bus not data bus.

MMU = Memory Management Unit

Pentium MMU is on same chip as CPU

Base-Limit Registers - Memory Protection + Address Translation

Several processes J,B,X.... in main memory

All processes have logical addresses starting at 0.

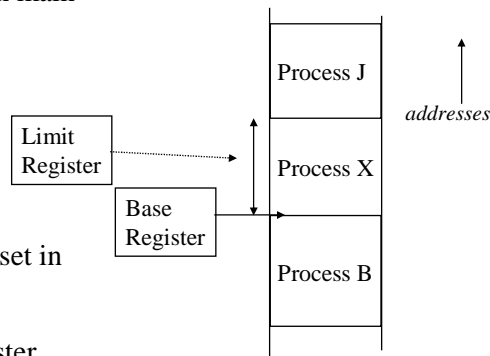
Process X is to run

Lowest physical address of X set in Base Register

Length of X set in Limit Register

Memory Protection: exception to CPU if address output by CPU \geq limit register contents

Translation: Logical Address + Base Register \longrightarrow Physical Address



All programs compiled to start at address 0, but program loaded anywhere in memory.

Before running a process, the Operating system sets the Base register with lowest address used by process, and Limit Register with length of process. Addresses used by process must be in range 0 to (length - 1)

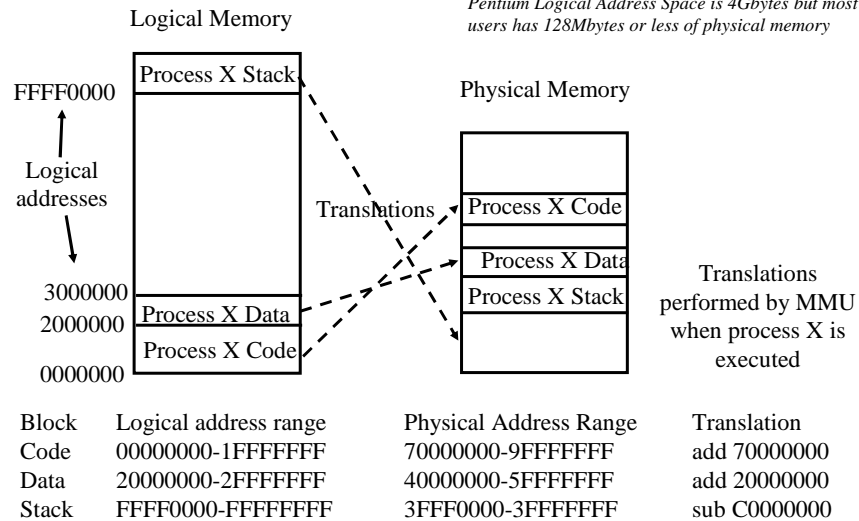
On each read and write cycle address output is checked against Limit register, and the process stopped through by an *exception* if the address is greater.

Otherwise the address from the CPU (the Logical Address) is *translated* to produce the Physical Address of the memory location to be accesses by adding the Base register contents:-

Logical Address + Base Register \rightarrow Physical Address

More Complex Address Translation Example

Address range of Logical Memory
is full address range of CPU.
Physical memory is usually smaller.
*Pentium Logical Address Space is 4Gbytes but most
users has 128Mbytes or less of physical memory*




Translation information created when program is loaded by O/S from disk into memory

Page-Based Memory Management - The standard memory management system

Process: Logical memory

Physical
memory



Program divided into equal size units:
pages.

Page size is always small and a power of 2:
256, 512,

Each page has separate translation to
Physical memory.

Translation information for program kept
in a "page table."

There is a page table for each process.

Page table can be very large

Translations only exist for Logical pages that are valid for process:
few, if any programs use all of Logical Address Space.

Attempt to access Logical address with no translation to physical memory
generates an exception to the CPU and the process is Terminated.

Page-based translation: Page Table

in general more complicated than previous slide

Translation information for program kept in page table

There is a page table for each process.

Part of Page Table for **a process** with 1000_{16} byte pages:

Logical page No	Logical Addresses	Physical Addresses
0	0000 - 0FFF	1A000 - 1AFFF
1	1000 - 1FFF	3000 - 3FFF
2	2000 - 2FFF	9000 - 9FFF
3	3000 - 3FFF	123000 - 123FFF
4	4000 - 4FFF	E8000 - E8FFF
5	5000 - 5FFF	4B000 - 4BFFF


From logical address, MMU gets page number and looks up table.

Physical addresses are allocated by O/S as program loaded into memory:
their values depend on what space is free at that time. O/S builds table.

Logical memory Space

of Program

conceptually divided into pages



Pentium:

Page size - 1000_{16}

Max No. of logical pages for
one program is 1000000_{16}

Advantages of Paging

- Can load program into available space in memory without need to coalesce free space into single block of free space.
- Can load several programs into memory at same time and run them even if programs' memory addresses are the same
[these are logical addresses which be translated to different physical addresses.
[each program has its own separate page table]
- compiler does not need to know memory addresses in which program will run:
compiler always use same memory addresses whatever the program
lets operating system create translation map on loading program.
- parts of logical address space unused by program have no mapping in table
- MMU generates an exception signal to CPU if access made to unused memory
and program aborted (bus error) *[part of protection system]*
- page table kept in memory - read from memory by MMU,
when process is scheduled to run

What's the use of address translation?

Better use of Memory

Can load programme into available space

Simpler compilation

compiler can compile all programs in same way with same logical addresses
[O/S sets up different translation for each programme loaded]

Protection of rest of memory

translations only for valid program addresses

generate exception to terminate programme if

invalid logical address sent to MMU

Prevents user program from accessing I/O interfaces and O/S code& data

[no valid translation provided to I/O addresses or O/S addresses]

Program is prevented from making changes to the page table:

no mapping to memory holding page table

O/S can access all addresses but User processes

can only access a restricted set of addresses

Virtual Memory an extension to paging

not all of program has to be in memory for program to run
keep some pages of program on disk and load as required to memory.

**Part of the page
table for a
process**


Logical page No	Logical Addresses	Physical Addresses
.		
96_{16}	96000 - 96FFF	19000 -19FFF
97_{16}	97000 - 97FFF	swapped to disk
98_{16}	98000 - 98FFF	not used by program
.		

accesses to page 96_{16} are translated by MMU and memory accessed

an access to page 98_{16} MMU generates an exception to CPU
O/S analyses exception and aborts program.

an access to page 97_{16} MMU generates an exception to CPU
O/S analyses exception and loads page from
disk into memory and process re-started.

Virtual Memory



Swap space is an area of a hard disk set aside for holding pages of memory swapped out to provide space for loading pages of another program

Virtual memory allows for better use of the available memory:

Large programs can be run on small memories: possibly more slowly.

More programs can have active pages loaded into memory.

Identification of O/S accesses from User Accesses

A CPU running address translation needs 2 or more modes of operation:-
e.g. a kernel or **supervisor mode** for the operating system
a **user mode** for application processes.


[Pentium has 4 modes or levels - with the Pentium everything is more complex]

Signal output by CPU indicates processor mode:

signal is used by external hardware to enable/disable access to parts of the system.

[MMU uses signal to distinguish between operating system and user process accesses.]

Must stop user mode process acquiring supervisor mode status.



This mode switch is made easy by the CPU.

This switch is only allowed via a system call instruction, which forces a switch to small number of predefined addresses in the O/S.

Process cannot switch the state and keep running.
The O/S is always entered on this switch.

Protection System